

Rahul Lotlikar

<https://github.com/rahul0776>

rahulujv@buffalo.edu

[linkedin.com/in/rahul-lotlikar-8a61141b2](https://www.linkedin.com/in/rahul-lotlikar-8a61141b2)

Buffalo, USA

Contact No. +1 (716) 344 6920

PROFESSIONAL SUMMARY

Software Engineer with 2+ years of experience in backend development, cloud integration, and data analytics, specializing in building scalable, fault-tolerant, and AI-driven systems. Proven track record in engineering end-to-end cloud pipelines at Minedco LiveOn, improving throughput by 40% and enabling real-time, production-grade AI solutions. Experienced in data engineering, statistical modeling, and workflow automation as a Research Assistant at SUNY Buffalo, and in delivering backend solutions and analytics platforms with Node.js, AWS, and PowerBI at Cybage Software. Currently pursuing an M.S. in Computer Science at SUNY Buffalo (Dec 2025) with a focus on Machine Learning, Deep Learning, and Computer Vision. Passionate about creating cloud-native, intelligent applications that solve real-world challenges at scale.

EDUCATION

University at Buffalo, SUNY Buffalo, NY
Master's Degree in Computer Science
Graduation Date: Dec 2025
Major Courses: Machine Learning, Deep Learning, Computer Vision

Rajiv Gandhi Institute of Technology Mumbai, India
Graduated with a Bachelor of Engineering in Information Technology
Graduation Date: May 2022

PROFESSIONAL EXPERIENCE

Minedco Holdings Inc. Leesburg, VA
Software Engineer Intern June 2025 - Aug 2025

- Built scalable backend pipelines with Azure Functions, Blob Storage, and serverless APIs to process millions of Facebook posts and images.
- Integrated LLMs for automated classification, summarization, and chapter creation, delivering production-grade AI solutions.
- Improved system throughput by 40% via parallelized data flows, asynchronous workloads, and FastAPI microservices.
- Developed cloud-native infrastructure (Azure, AWS) and CI/CD pipelines for reliable, distributed deployments.

Cybage Software Pvt. Ltd Pune, India
Software Engineer Jul 2022 - Jun 2024

- Designed and implemented scalable backend solutions using Node.js, AWS Lambda, and DynamoDB, improving system performance by 30% and reducing operational costs by 20%.
- Developed and maintained data pipelines and storage systems with AWS DynamoDB, ensuring 99.9% uptime and enhanced data security compliance.
- Collaborated with GroupM to deliver advanced data analytics and PowerBI dashboards, accelerating decision-making by 25% and optimizing resource allocation.
- Integrated Google Ads and Meta Ads data to improve visibility across 50+ digital campaigns, contributing to a 15% increase in overall marketing efficiency.
- Resolved 80% of critical backend issues and deployed fault-tolerant systems serving over 100,000 users, improving stability and user experience across multiple client platforms.

TECHNICAL SKILLS

Programming Languages: HTML, Java, TypeScript, JavaScript, Python

Frameworks and Libraries: Node.js, Angular, FastAPI, TensorFlow, PyTorch, Pandas

Cloud and Analytics Tools: AWS (Lambda, DynamoDB, S3), Azure Functions, CI/CD, Serverless Architectures, Containerization (Docker, Kubernetes)

Databases: MySQL, MongoDB, DynamoDB

Data Engineering & Analytics: Distributed Systems, APIs, Microservices, Data Pipelines, PowerBI, Tableau, Pandas, Excel

Version Control Systems: Git, GitHub

PROJECTS

AI Document Search (RAG Chatbot) | *React, FastAPI, LangChain, Pinecone/FAISS, Docker, Vercel*

- Built a full-stack "Chat with PDFs" system enabling users to upload documents and query them via a retrieval-augmented generation (RAG) pipeline with semantic search, LLM orchestration, and citation-backed answers.
- Engineered end-to-end data pipelines: text extraction, embedding generation (OpenAI), vector search (FAISS/Pinecone), and token-efficient answer generation with LangChain.
- Deployed frontend on Vercel and backend as a Dockerized FastAPI service with CI/CD, monitoring, and observability. Achieved < 800ms retrieval latency, < 2.5s first-token response, and 99.9% processing reliability.
- Implemented advanced features including streaming token responses, reranking, caching (30–50% cost savings), recall@5 ≥ 0.85, and hallucination rate < 5%, showcasing reliability, scalability, and MLOps best practices.

AI-Powered Marketing Campaign Optimizer *Technologies: Python, Streamlit, Scikit-learn, Pandas, Matplotlib, Joblib, Cython, Pillow*

- Built and deployed a Streamlit web app that predicts customer responses with 92% accuracy, using machine learning pipelines, feature engineering, and real-time inference.
- Delivered interactive dashboards with feature importance visualization on Streamlit Cloud, enabling marketers to optimize campaigns and resource allocation.